

Homework-2-2

March 1, 2018

1 Clustering Yelp Restaurants

Part 2: To be completed INDIVIDUALLY and due on March 3 at 7pm.

In this assignment, we will continue to work with the [Yelp dataset](#) that we used in Homework 2-1.

We will continue to try to find culinary **districts** in Las Vegas. As a reminder from last time, these are characterized by **closeness** and **similarity** of restaurants. Use the "longitude" and "latitude" to cluster closeness.

However, in this analysis we will not use the Yelp-supplied "categories" to cluster for similarity as we did in Part 1. Instead, we will cluster the reviews themselves, extracting categories in an unsupervised fashion.

Specifically, you are to use Latent Semantic Analysis (LSA) on the Yelp reviews to cluster restaurants based off on their reviews. As a reminder, LSA consists of using PCA applied to the document-term matrix.

Now, your feature vectors will contain latitude, longitude and the most relatively important review terms.

You will apply PCA 3 times. Each time, you will take into account the first k reviews per business, where $k = \{10, 100, 1000\}$. Many businesses will have less than k , or even no reviews. In this case, simply assign to the business the maximum number of reviews it has. **(4 pts)**

```
In [1]: import json
```

```
%matplotlib inline
%config InlineBackend.figure_format='retina'
# import libraries
import numpy as np
import matplotlib as mp
import pandas as pd
import matplotlib.pyplot as plt
import pandas as pd
#import slideUtilities as sl
#import laUtilities as ut
from importlib import reload
from datetime import datetime
from IPython.display import Image
from IPython.display import display_html
from IPython.display import display
```

```

from IPython.display import Math
from IPython.display import Latex
from IPython.display import HTML
import sklearn.datasets as sk_data
import sklearn.metrics as metrics

```

```

#import matplotlib as mpl
import seaborn as sns
print('')

```

```

In [2]: restaurant = {}
        with open('dataset/business.json') as bus:
            for line in bus:
                dic = json.loads(line)
                if dic["city"]=="Las Vegas":
                    if "Restaurants" in dic["categories"]:
                        #
                        kk[dic["business_id"]] = 0
                        restaurant[dic["business_id"]] = {"latitude":dic["latitude"], "longitude":dic["longitude"]}

df = pd.DataFrame.from_dict(restaurant, orient='index')
df.reset_index(level=0, inplace=True)
df.head()

```

```

Out[2]:
          index  latitude  longitude
0  --9e10NYQuAa-CB_Rrw7Tw  36.123183 -115.169190
1  --q7kSBRb0vWC81SkXFBYA  36.016693 -115.173115
2  -153AjTW5luZPK4omEujWA  36.103001 -115.173516
3  -1m9o3vGRA8IBPNvNqKLmA  36.104330 -115.175593
4  -1vfRrlnNnNJ5bo0VghMPA  36.281295 -115.286737

```

```

In [3]: latitude = df.latitude
        longitude = df.longitude
        location = np.array(list(zip(latitude, longitude)))

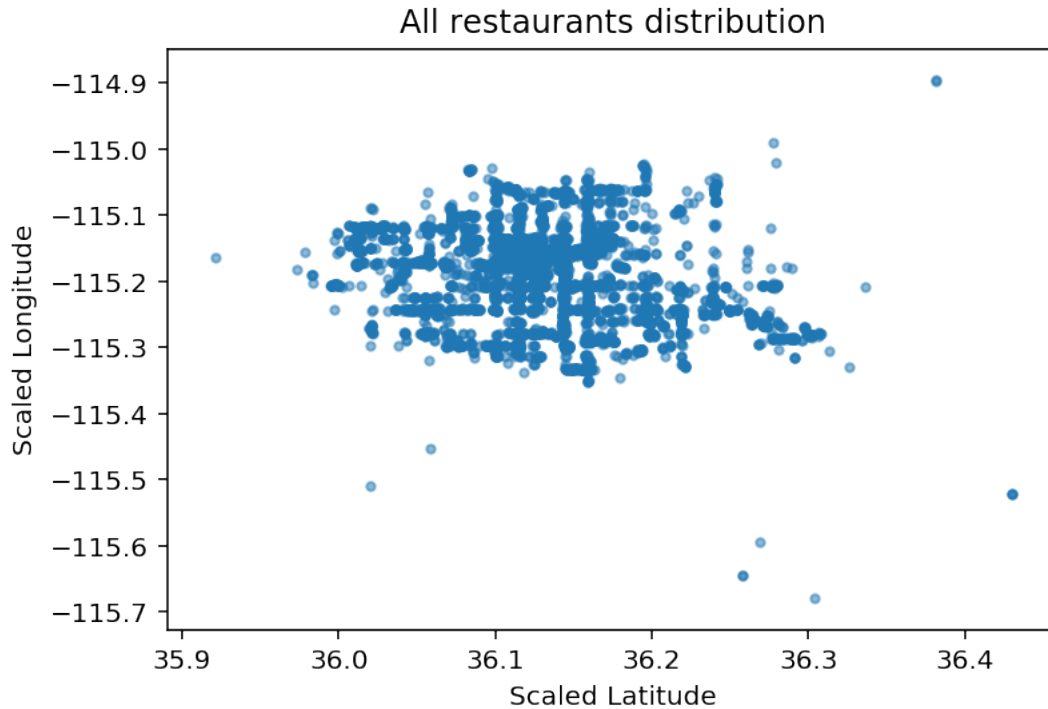
        plt.scatter(location[:, 0],location[:, 1], s=10, alpha = 0.5)
        plt.title('All restaurants distribution')
        plt.xlabel('Scaled Latitude')
        plt.ylabel('Scaled Longitude')

```

```

Out[3]: Text(0,0.5,'Scaled Longitude')

```



```
In [4]: from collections import defaultdict

review_dic = {}

with open('dataset/review.json') as f:
    for line in f:
        dic1 = json.loads(line)
        if dic1["business_id"] in restaurant:
            if dic1["business_id"] in review_dic:
                review_dic[dic1["business_id"]].append((dic1["text"]))
            else:
                review_dic[dic1["business_id"]] = [dic1["text"]]

In [5]: print(len(review_dic))
        print(len(review_dic['zpoZ6WyQUYff18-z4ZU1mA']))

review10_dic = {}
review10_str = {}
review100_dic = {}
review100_str = {}
review1000_dic = {}
review1000_str = {}
for key in review_dic:
```

```

m10 = min(len(review_dic[key]), 10)
m100 = min(len(review_dic[key]), 100)
m1000 = min(len(review_dic[key]), 1000)

review10_dic[key] = review_dic[key][:m10]
review10_str[key] = ' '.join(review10_dic[key])
review100_dic[key] = review_dic[key][:m100]
review100_str[key] = ' '.join(review100_dic[key])
review1000_dic[key] = review_dic[key][:m1000]
review1000_str[key] = ' '.join(review1000_dic[key])

```

5899

547

In [6]: *# create separate lists of business_id, latitude, longitude and review, with the same*

```

business_id = []
latitude = []
longitude = []
review10 = []
review100 = []
review1000 = []

for key, val in restaurant.items():
    business_id.append(key)
    latitude.append(val.get('latitude'))
    longitude.append(val.get('longitude'))
    review10.append(review10_str[key])
    review100.append(review100_str[key])
    review1000.append(review1000_str[key])

```

```

latitude = np.array(latitude)
longitude = np.array(longitude)

```

In [7]: *# strip digits from all reviews*

```

from string import digits
remove_digits = str.maketrans("", "", digits)
review10 = [rev.translate(remove_digits) for rev in review10]
review100 = [rev.translate(remove_digits) for rev in review100]
review1000 = [rev.translate(remove_digits) for rev in review1000]

print(len(review10))

```

5899

In [8]: `from nltk.stem.snowball import SnowballStemmer`
`from nltk.tokenize import word_tokenize`

```

from nltk.tokenize import sent_tokenize

review10 = [" ".join(SnowballStemmer("english", ignore_stopwords=True).stem(word)
                    for sent in sent_tokenize(message)
                    for word in word_tokenize(sent))
            for message in review10]

In [9]: review100 = [" ".join(SnowballStemmer("english", ignore_stopwords=True).stem(word)
                             for sent in sent_tokenize(message)
                             for word in word_tokenize(sent))
                    for message in review100]

```

KeyboardInterrupt

Traceback (most recent call last)

```

<ipython-input-9-403ecd5d8f48> in <module>()
      2         for sent in sent_tokenize(message)
      3         for word in word_tokenize(sent))
----> 4         for message in review100]
      5

<ipython-input-9-403ecd5d8f48> in <listcomp>(.0)
      2         for sent in sent_tokenize(message)
      3         for word in word_tokenize(sent))
----> 4         for message in review100]
      5

<ipython-input-9-403ecd5d8f48> in <genexpr>(.0)
      1 review100 = [" ".join(SnowballStemmer("english", ignore_stopwords=True).stem(word)
      2         for sent in sent_tokenize(message)
----> 3         for word in word_tokenize(sent))
      4         for message in review100]
      5

/Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages/nltk/stemmer.py
    96         raise ValueError("The language '{0}' is not supported.".format(language))
    97     stemmerclass = globals()[language.capitalize() + "Stemmer"]
--> 98     self.stemmer = stemmerclass(ignore_stopwords)
    99     self.stem = self.stemmer.stem
   100     self.stopwords = self.stemmer.stopwords

```

```

/Library/Frameworks/Python.framework/Versions/3.6/lib/python3.6/site-packages/nltk/stem
130         try:
131             for word in stopwords.words(language):
--> 132                 self.stopwords.add(word)
133         except IOError:
134             raise ValueError("{!r} has no list of stopwords. Please set"

```

KeyboardInterrupt:

```

In [ ]: review1000 = [" ".join(SnowballStemmer("english", ignore_stopwords=True).stem(word)
        for sent in sent_tokenize(message)
        for word in word_tokenize(sent))
        for message in review1000]

```

```

In [185]: from sklearn.feature_extraction.text import TfidfVectorizer
        # min_df=4, max_df=0.8
        # max_df=0.18, min_df=0.009
        vectorizer = TfidfVectorizer(min_df=0.009, max_df=0.18, stop_words='english', sublinea
        DataReview10 = vectorizer.fit_transform(review10)
        terms10 = vectorizer.get_feature_names()

        print(type(DataReview10), DataReview10.shape)
        # terms10 = vectorizer.get_feature_names()

```

```

<class 'scipy.sparse.csr.csr_matrix'> (5899, 2591)

```

```

In [194]: DataReview100 = vectorizer.fit_transform(review100)
        terms100 = vectorizer.get_feature_names()

        print(type(DataReview100), DataReview100.shape)

```

```

<class 'scipy.sparse.csr.csr_matrix'> (5899, 7982)

```

```

In [211]: DataReview1000 = vectorizer.fit_transform(review1000)
        terms1000 = vectorizer.get_feature_names()

        print(type(DataReview1000), DataReview1000.shape)

```

```

<class 'scipy.sparse.csr.csr_matrix'> (5899, 12465)

```

```

In [23]: from sklearn.decomposition import TruncatedSVD
        from sklearn.preprocessing import Normalizer
        # K = 10
        # apply LSA on reviews


```

```

variance = np.zeros(10)
ii = 0
for kk in np.linspace(1,501, 10, dtype='int16'):
    lsa = TruncatedSVD(kk, algorithm = 'arpack')
    DataReview10_lsa = lsa.fit(DataReview10)
    variance[ii] = DataReview10_lsa.explained_variance_ratio_.mean()
    ii += 1

plt.plot(np.linspace(1,501, 10, dtype='int16'),variance)
plt.xlabel('Number of components')
dummy = plt.ylabel('variance')
plt.show()

```



```

In [186]: from sklearn.decomposition import TruncatedSVD
          from sklearn.preprocessing import Normalizer
          # select 100 to apply to SVD
          lsa = TruncatedSVD(100, algorithm = 'arpack')
          review_lsa10 = lsa.fit_transform(DataReview10)
          review_lsa10 = Normalizer(copy=False).fit_transform(review_lsa10)

In [195]: lsa = TruncatedSVD(100, algorithm = 'arpack')
          review_lsa100 = lsa.fit_transform(DataReview100)
          review_lsa100 = Normalizer(copy=False).fit_transform(review_lsa100)

In [212]: lsa = TruncatedSVD(100, algorithm = 'arpack')
          review_lsa1000 = lsa.fit_transform(DataReview1000)
          review_lsa1000 = Normalizer(copy=False).fit_transform(review_lsa1000)

```

```
In [74]: location = 1000 * Normalizer(copy=False).fit_transform(np.column_stack((latitude, longitude)))
feature10 = np.concatenate((review_lsa10, location), axis=1)

In [196]: feature100 = np.concatenate((review_lsa100, location), axis=1)

In [213]: feature1000 = np.concatenate((review_lsa1000, location), axis=1)
```

Find clusters using the 3 different techniques we discussed in class: k-means++, hierarchical, and GMM for **each** of the 3 feature vectors per business (remember you created feature vectors based on $k = \{10, 100, 1000\}$). Visualize the clusters by plotting the longitude/latitude of the restaurants in a scatter plot and label each cluster.

Note that to label each cluster, you will need to think about how to extract labels from the LSA results. (4 pts)

2 K=10: first 10 reviews per business


3 K-Means++ clusters

```
In [187]: from sklearn.cluster import KMeans
          from sklearn import metrics
          import numpy as np
          import seaborn as sns; sns.set() # for plot styling

          max_clusters = 15
          s = np.zeros(max_clusters+1)

          for k in range(2, max_clusters+1):
              kmeans = KMeans(init='k-means++', n_clusters=k, n_init=10)
              kmeans.fit_predict(feature10)
              labels = kmeans.fit_predict(feature10)
              s[k] = metrics.silhouette_score(feature10, labels, metric='euclidean')

          plt.plot(range(2, len(s)), s[2:])
          plt.xlabel('Number of clusters')
          plt.ylabel('Kmeans Silhouette Score')
          plt.show()
```

```
In [76]: # select k=5 to cluster
k_K10 = 5

kmeans = KMeans(init='k-means++', n_clusters=k_K10, n_init=10)
K_labels10 = kmeans.fit_predict(feature10)
K_centroids10 = kmeans.cluster_centers_

K_dic10 = {}
for i in K_labels10:
    if i in K_dic10:
        K_dic10[i] += 1
    else:
        K_dic10[i] = 1

# print(K_dic10)
for i in range(k_K10):
    print('There are {} restautrants in cluster {}'.format(K_dic10[i], i))
```

There are 795 restautrants in cluster 0.
 There are 1070 restautrants in cluster 1.
 There are 481 restautrants in cluster 2.
 There are 1184 restautrants in cluster 3.
 There are 2369 restautrants in cluster 4.

```
In [83]: K_centroids10 = kmeans.cluster_centers_
K_ori_centers10 = lsa.inverse_transform(K_centroids10[:, :100])
K_order_centroids10 = K_ori_centers10.argsort()[:, :-1]

print("Top words in each cluster by using K-means method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_K10):
    print("Top 50 words in Cluster {}: ".format(i))
    for ind in K_order_centroids10[i, :50]:
        # print('Top category is {}'.format(terms10[ind]))
        print(' %s' % terms10[ind], end='')
    print("")
    print("")
```

Top words in each cluster by using K-means method:

Top 50 words in Cluster 0:

noodl sushi chines thai bowl tea asian korean broth pho japanes curri tofu chef authent pad sp

Top 50 words in Cluster 1:

wing taco bartend garlic bowl beer bbq deliveri game crust sushi ingredi sub play burrito read

Top 50 words in Cluster 2:

taco mexican salsa burrito asada carn chip tortilla enchilada guacamol authent roberto pastor

Top 50 words in Cluster 3:

wing downtown taco deliveri beer bacon cashier bbq window phone card deliv subway job readi c

Top 50 words in Cluster 4:

hotel coffe beer casino bartend bacon waiter wing wine italian expans tomato buffet pasta caf

So From top 50 words in each cluster, we can selects several key words as their labels.

Cluster 0 : noodles, sushi, Chinese

Cluster 1 : wings, taco, BBQ

Cluster 2 : Mexican, salsa, burrito

Cluster 3 : downtown, beer, bacon

Cluster 4 : coffee, wine, Italian

```
In [86]: K_cluster10_top = [['noodles, sushi, Chinese'], ['wings, taco, BBQ'], ['Mexican, salsa, burrito'], ['downtown, beer, bacon'], ['coffee, wine, Italian']]


c = ['r', 'g', 'blue', 'yellow', 'purple', 'darkseagreen', 'dimgray', 'pink', 'brown', 'brown']

plt.figure(figsize=(16,10))

for i in range(k_K10):
    plt.scatter(latitude[K_labels10==i], longitude[K_labels10==i], s=25, color=c[i], alpha=0.5)
#     points = np.array([location[j] for j in range(len(location)) if K_labels[j] == i])
#     plt.scatter(points[:, 0], points[:, 1], s=10, c=c[i], alpha = 0.5)

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('Kmeans Clusters of top10 reviews in each restaurants.')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```

Out[86]: Text(0,0.5,'Longitude')



4 Hierarchical clusters

```
In [89]: import scipy.cluster
import scipy.cluster.hierarchy as hierarchy
import scipy.spatial.distance


H_Z10 = hierarchy.linkage(feature10, method='ward', metric='euclidean')
```

```

s = np.zeros(max_clusters+1)
for k in range(2,max_clusters+1):
    clusters = hierarchy.fcluster(H_Z10, k, criterion='maxclust')
    s[k] = metrics.silhouette_score(feature10,clusters,metric='euclidean')
plt.plot(range(2,len(s)),s[2:])
plt.xlabel('Number of clusters')
plt.ylabel('Hierarchy Silhouette Score')

```

Out[89]: Text(0,0.5,'Hierarchy Silhouette Score')



```

In [90]: k_H10 = 5
clusters10 = hierarchy.fcluster(H_Z10, k_H10, criterion='maxclust')

H_dic10 = {}
for i in clusters10:
    if i in H_dic10:
        H_dic10[i] += 1
    else:
        H_dic10[i] = 1


for i in range(1,k_H10+1):
    print('There are {} restauetrants in cluster {}'.format(H_dic10[i], i))

```

There are 1151 restauetrants in cluster 1.
There are 489 restauetrants in cluster 2.

There are 3290 restaurants in cluster 3.
 There are 330 restaurants in cluster 4.
 There are 639 restaurants in cluster 5.

```
In [91]: plt.figure(figsize=(14,4))
         H_figure = hierarchy.dendrogram(H_Z10, truncate_mode='level', show_leaf_counts=True)
```



```
In [103]: # Calculate the centroids in each cluster
          cluster_H10 = {}
          sum_H10 = []
          H_centroids10 = []
          sum_HH10 = np.zeros(102)
          for i in range(k_H10):
              cluster_H10[i] = np.array([feature10[j] for j in range(len(location)) if cluster
              for k in range(len(cluster_H10[i])):
                  sum_HH10 = list(map(lambda x: x[0]+x[1], zip(sum_HH10, cluster_H10[i][k])))
              sum_H10.append(sum_HH10)
              center = [(sum_HH10[j]/len(cluster_H10[i])) for j in range(len(sum_HH10))]
              H_centroids10.append(center)

          H_centroids10 = np.array(H_centroids10)
          H_ori_centers10 = lsa.inverse_transform(H_centroids10[:, :100])
          H_order_centroids10 = H_ori_centers10.argsort()[:, :-1]

          print("Top words in each cluster by using Hierarchical method:")
          print('')

          # print out the top 50 words with largest weight in each cluster
          for i in range(k_H10):
              print("Top 50 words in Cluster {}: ".format(i+1))
              for ind in H_order_centroids10[i, :50]:
                  print(' %s' % terms10[ind], end='')
              print("")
              print("")
```

Top words in each cluster by using Hierarchical method:

Top 50 words in Cluster 1:

taco mexican burrito salsa asada carn chip bartend beer tortilla truck downtown el wing bacon

Top 50 words in Cluster 2:

chines taco noodl thai mexican burrito authent salsa asada deliveri curri bbq sour chip carn

Top 50 words in Cluster 3:

beer taco wing bartend noodl chines bbq waiter coffe bowl bacon garlic sushi hotel deliveri

Top 50 words in Cluster 4:

taco beer mexican chip burrito wing bartend waiter noodl bowl chines authent bbq salsa hotel

Top 50 words in Cluster 5:

taco wing beer mexican chip burrito bowl bartend garlic coffe bacon bbq waiter sushi deliveri

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : taco, Mexican, burrito

Cluster 1 : Chinese, noodles, Thai

Cluster 2 : beer, wings, BBQ

Cluster 3 : Mexican, bowl, sushi


Cluster 4 : taco, chip, coffee

```
In [108]: H_cluster10_top = [['taco, Mexican, burrito'], ['Chinese, noodles, Thai'], ['beer, wing
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_H10):
    plt.scatter(latitude[clusters10==i+1], longitude[clusters10==i+1], s=25, color=c[i])

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('Hierachy Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')


Out[108]: Text(0,0.5, 'Longitude')
```



5 GMM clusters

```
In [189]: s = np.zeros(max_clusters+1)
          for k in range(2,max_clusters+1):
              gmm = mixture.GaussianMixture(n_components=k, covariance_type='full')
              gmm.fit(feature10)
              y_pred = gmm.predict(feature10)
              s[k] = metrics.silhouette_score(feature10,y_pred,metric='euclidean')
          plt.plot(range(2,len(s)),s[2:])
          plt.xlabel('Number of clusters')
          plt.ylabel('GMM Silhouette Score')
```

```
Out[189]: Text(0,0.5,'GMM Silhouette Score')
```



```
In [193]: k_G10 = 5
          gmm = mixture.GaussianMixture(n_components=k_G10, covariance_type='tied')
          gmm.fit(feature10)
          y_pred10 = gmm.predict(feature10)

          G_dic10 = {}
          for i in y_pred10:
              if i in G_dic10:
                  G_dic10[i] += 1
              else:
                  G_dic10[i] = 1

          for i in range(k_G10):
              print('There are {} restautrants in cluster {}'.format(G_dic10[i], i))
```

```
There are 3293 restautrants in cluster 0.
There are 537 restautrants in cluster 1.
There are 529 restautrants in cluster 2.
There are 511 restautrants in cluster 3.
There are 1029 restautrants in cluster 4.
```

```
In [191]: # Calculate the centroids in each cluster
          cluster_G10 = {}
          sum_G10 = []
```



```

G_centroids10 = []
sum_GG10 = np.zeros(102)
for i in range(k_G10):
    cluster_G10[i] = np.array([feature10[j] for j in range(len(location)) if y_pred1
    for k in range(len(cluster_G10[i])):
        sum_GG10 = list(map(lambda x: x[0]+x[1], zip(sum_GG10, cluster_G10[i][k])))
    sum_G10.append(sum_GG10)
    center = [(sum_GG10[j]/len(cluster_G10[i])) for j in range(len(sum_GG10))]
    G_centroids10.append(center)

G_centroids10 = np.array(G_centroids10)
G_ori_centers10 = lsa.inverse_transform(G_centroids10[:, :100])
G_order_centroids10 = G_ori_centers10.argsort()[:, :-1]

print("Top words in each cluster by using GMM method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_G10):
    print("Top 50 words in Cluster {}: ".format(i))
    for ind in G_order_centroids10[i, :50]:
        print(' %s' % terms10[ind], end='')
    print("")
    print("")

```

Top words in each cluster by using GMM method:

Top 50 words in Cluster 0:

sushi noodl thai curri chines bowl korean pho tea japanes broth ayc chef pad tempura tofu asi

Top 50 words in Cluster 1:

sushi noodl subway thai sub bowl curri chines tea korean pho tuna japanes broth chef ayc pad a

Top 50 words in Cluster 2:

wing beer bartend garlic sushi bbq bowl noodl deliveri chines waiter chef bacon coffe tea rib

Top 50 words in Cluster 3:

wing beer coffe garlic bbq bartend sushi bacon bowl noodl deliveri chines tea waiter chef rib

Top 50 words in Cluster 4:

taco wing beer mexican chip burrito bowl bartend garlic coffe bacon bbq waiter sushi deliveri

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : sushi, wine, lobster

Cluster 1 : taco, Mexican, salsa

Cluster 2 : beer, noodle, Chinese

Cluster 3 : taco, chip, coffee


Cluster 4 : wings, bacon, BBQ

```
In [118]: G_cluster10_top = [['sushi, wine, lobster'], ['taco, Mexican, salsa'], ['beer, noodle,
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_G10):
    plt.scatter(latitude[y_pred10==i], longitude[y_pred10==i], s=25, color=c[i], label=i)

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('GMM Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```

Out[118]: Text(0,0.5, 'Longitude')



6 K=100: first 100 reviews per business

7 K-Means++ clusters


```
In [197]: max_clusters = 15
s = np.zeros(max_clusters+1)
```

```

for k in range(2,max_clusters+1):
    kmeans = KMeans(init='k-means++', n_clusters=k, n_init=10)
    kmeans.fit_predict(feature100)
    labels = kmeans.fit_predict(feature100)
    s[k] = metrics.silhouette_score(feature100,labels,metric='euclidean')

plt.plot(range(2,len(s)),s[2:])
plt.xlabel('Number of clusters')
plt.ylabel('Kmeans Silhouette Score')
plt.show()

```



```

In [198]: # select k=5 to cluster
k_K100 = 5

kmeans = KMeans(init='k-means++', n_clusters=k_K100, n_init=10)
K_labels100 = kmeans.fit_predict(feature100)
K_centroids100 = kmeans.cluster_centers_

K_dic100 = {}
for i in K_labels100:
    if i in K_dic100:
        K_dic100[i] += 1
    else:
        K_dic100[i] = 1

```

```

# print(K_dic10)
for i in range(k_K100):
    print('There are {} restautrants in cluster {}. '.format(K_dic100[i], i))

```

There are 458 restautrants in cluster 0.
 There are 1207 restautrants in cluster 1.
 There are 1058 restautrants in cluster 2.
 There are 836 restautrants in cluster 3.
 There are 2340 restautrants in cluster 4.

```

In [199]: K_centroids100 = kmeans.cluster_centers_
K_ori_centers100 = lsa.inverse_transform(K_centroids100[:, :100])
K_order_centroids100 = K_ori_centers100.argsort()[:, :-1]

print("Top words in each cluster by using K-means method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_K100):
    print("Top 50 words in Cluster {}: ".format(i))
    for ind in K_order_centroids100[i, :50]:
        # print('Top category is {}'.format(terms10[ind]))
        print(' %s' % terms100[ind], end='')
    print("")
    print("")

```

Top words in each cluster by using K-means method:

Top 50 words in Cluster 0:

asada carne tortillas burritos guacamole pastor carnitas enchiladas el tortilla nachos salsas

Top 50 words in Cluster 1:

downtown subway pizzas fremont mac pepperoni truck refund pie ribs bell fingers mcdonald subs

Top 50 words in Cluster 2:

subway pizzas pepperoni mac bartenders fingers sushi buffalo teriyaki pie philly nachos bell

Top 50 words in Cluster 3:

sushi noodle tofu thai japanese broth curry korean chow mein tempura soy pad pho miso chinato

Top 50 words in Cluster 4:

court pizzas mac subway pepperoni caesar lobster pie york pool bartenders mashed dogs lamb ha

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : burritos, nachos, Mexican

Cluster 1 : pizzas, pie, burgers


Cluster 2 : subway, fingers, meatballs

Cluster 3 : sushi, Japanese, Chinese

Cluster 4 : lobster, brunch, pancakes

```
In [200]: K_cluster100_top = [['burritos, nachos, Mexican'], [' pizzas, pie, burgers'], ['subway  
  
c = ['r', 'g', 'blue', 'yellow', 'purple', 'darkseagreen', 'dimgray', 'pink', 'brown'  
  
plt.figure(figsize=(16,10))  
  
for i in range(k_K100):  
    plt.scatter(latitude[K_labels100==i], longitude[K_labels100==i], s=25, color=c[i])  
  
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})  
plt.title('Kmeans Clusters of top100 reviews in each restaurants.')  
plt.xlabel('Latitude')  
plt.ylabel('Longitude')
```

```
Out[200]: Text(0,0.5,'Longitude')
```




8 Hierarchical clusters

```
In [201]: H_Z100 = hierarchy.linkage(feature100, method='ward', metric='euclidean')

s = np.zeros(max_clusters+1)
for k in range(2,max_clusters+1):
    clusters = hierarchy.fcluster(H_Z100, k, criterion='maxclust')
    s[k] = metrics.silhouette_score(feature100,clusters,metric='euclidean')
plt.plot(range(2,len(s)),s[2:])
plt.xlabel('Number of clusters')
plt.ylabel('Hierarchy Silhouette Score')

Out[201]: Text(0,0.5,'Hierarchy Silhouette Score')
```



```
In [204]: k_H100 = 5
clusters100 = hierarchy.fcluster(H_Z100, k_H100, criterion='maxclust')

H_dic100 = {}
for i in clusters100:
    if i in H_dic100:
        H_dic100[i] += 1
    else:
        H_dic100[i] = 1

for i in range(1,k_H100+1):
    print('There are {} restaurants in cluster {}'.format(H_dic100[i], i))
```

There are 403 restaurants in cluster 1.
 There are 463 restaurants in cluster 2.
 There are 835 restaurants in cluster 3.
 There are 951 restaurants in cluster 4.
 There are 3247 restaurants in cluster 5.

```
In [205]: # Calculate the centroids in each cluster
cluster_H100 = {}
sum_H100 = []
H_centroids100 = []
sum_HH100 = np.zeros(102)
for i in range(k_H100):
    cluster_H100[i] = np.array([feature100[j] for j in range(len(location)) if cluster_H100[i][j] == feature100[j]])
    for k in range(len(cluster_H100[i])):
        sum_HH100 = list(map(lambda x: x[0]+x[1], zip(sum_HH100, cluster_H100[i][k])))
    sum_H100.append(sum_HH100)
    center = [(sum_HH100[j]/len(cluster_H100[i])) for j in range(len(sum_HH100))]
    H_centroids100.append(center)

H_centroids100 = np.array(H_centroids100)
H_ori_centers100 = lsa.inverse_transform(H_centroids100[:, :100])
H_order_centroids100 = H_ori_centers100.argsort()[:, :-1]

print("Top words in each cluster by using Hierarchical method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_H100):
    print("Top 50 words in Cluster {}: ".format(i+1))
    for ind in H_order_centroids100[i, :50]:
        print(' %s' % terms100[ind], end='')
    print("")
    print("")
```

Top words in each cluster by using Hierarchical method:

Top 50 words in Cluster 1:

asada carne el pastor tortillas burritos carnitas horchata guacamole al salsas enchiladas tor

Top 50 words in Cluster 2:

asada carne el tortillas burritos pastor truck guacamole tortilla nachos carnitas horchata al

Top 50 words in Cluster 3:

asada carne burritos tortillas guacamole nachos tortilla el carnitas quesadilla enchiladas tr

Top 50 words in Cluster 4:

asada carne sushi thai burritos noodle tortillas tofu guacamole curry nachos tortilla teriyaki

Top 50 words in Cluster 5:

sushi subway pizzas asada nachos mac carne pepperoni ribs court thai lobster burritos noodle

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 1 : burritos, Mexican, nachos

Cluster 2 : carne, pizzas, Spanish

Cluster 3 : fingers, sushi, downtown

Cluster 4 : Thai, tofu, Japanese

Cluster 5 : subway, noodles, pancakes

```
In [206]: H_cluster100_top = [['burritos, Mexican, nachos'], ['carne, pizzas, Spanish'], ['fingers, sushi, downtown'],
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_H100):
    plt.scatter(latitude[clusters100==i+1], longitude[clusters100==i+1], s=25, color=c[i])

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('Hierachy Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```


Out[206]: Text(0,0.5,'Longitude')



9 GMM clusters

```
In [207]: s = np.zeros(max_clusters+1)
          for k in range(2,max_clusters+1):
              gmm = mixture.GaussianMixture(n_components=k, covariance_type='full')
              gmm.fit(feature100)
              y_pred = gmm.predict(feature100)
              s[k] = metrics.silhouette_score(feature100,y_pred,metric='euclidean')
          plt.plot(range(2,len(s)),s[2:])
          plt.xlabel('Number of clusters')
          plt.ylabel('GMM Silhouette Score')
```

```
Out[207]: Text(0,0.5,'GMM Silhouette Score')
```



```
In [208]: k_G100 = 5
          gmm = mixture.GaussianMixture(n_components=k_G100, covariance_type='tied')
          gmm.fit(feature100)
          y_pred100 = gmm.predict(feature100)

          G_dic100 = {}
          for i in y_pred100:
              if i in G_dic100:
```

```

        G_dic100[i] += 1
    else:
        G_dic100[i] = 1

    for i in range(k_G100):
        print('There are {} restautrants in cluster {}'.format(G_dic100[i], i))

```

There are 451 restautrants in cluster 0.
 There are 3029 restautrants in cluster 1.
 There are 584 restautrants in cluster 2.
 There are 653 restautrants in cluster 3.
 There are 1182 restautrants in cluster 4.

```

In [209]: # Calculate the centroids in each cluster
cluster_G100 = {}
sum_G100 = []
G_centroids100 = []
sum_GG100 = np.zeros(102)
for i in range(k_G100):
    cluster_G100[i] = np.array([feature100[j] for j in range(len(location)) if y_pre == i])
    for k in range(len(cluster_G100[i])):
        sum_GG100 = list(map(lambda x: x[0]+x[1], zip(sum_GG100, cluster_G100[i][k])))
    sum_G100.append(sum_GG100)
    center = [(sum_GG100[j]/len(cluster_G100[i])) for j in range(len(sum_GG100))]
    G_centroids100.append(center)

G_centroids100 = np.array(G_centroids100)
G_ori_centers100 = lsa.inverse_transform(G_centroids100[:, :100])
G_order_centroids100 = G_ori_centers100.argsort()[:, :-1]

print("Top words in each cluster by using GMM method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_G100):
    print("Top 50 words in Cluster {}: ".format(i))
    for ind in G_order_centroids100[i, :50]:
        print(' %s' % terms100[ind], end='')
    print("")
    print("")

```

Top words in each cluster by using GMM method:

Top 50 words in Cluster 0:

pizzas pepperoni hut dough driver oven pie knots topping calzone pizzeria mozzarella papa dom

Top 50 words in Cluster 1:

subway pizzas pepperoni court pie mac mcdonald refund driver dough dogs oven subs hash buffalo

Top 50 words in Cluster 2:

subway asada pizzas carne pepperoni mac court pie nachos dogs refund el guacamole ribs dough m

Top 50 words in Cluster 3:

subway asada pizzas nachos carne pepperoni burritos mac court tortilla guacamole pie refund b

Top 50 words in Cluster 4:

sushi subway pizzas asada nachos mac carne pepperoni ribs court thai lobster burritos noodle p

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : pizzas, pie, meatballs

Cluster 1 : subway, pancakes, burger

Cluster 2 : nachos, Mexican, burrito

Cluster 3 : carne, fingers, Chinese

Cluster 4 : sushi, noodles, curry

```
In [210]: G_cluster100_top = [['pizzas', 'pie', 'meatballs'], ['subway', 'pancakes', 'burger'], ['nachos',
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_G100):
    plt.scatter(latitude[y_pred100==i], longitude[y_pred100==i], s=25, color=c[i], label=i)

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('GMM Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')

Out[210]: Text(0,0.5, 'Longitude')
```




10 K=1000: first 1000 reviews per business

11 K-Means++ clusters

```
In [166]: max_clusters = 15
          s = np.zeros(max_clusters+1)

          for k in range(2,max_clusters+1):
              kmeans = KMeans(init='k-means++', n_clusters=k, n_init=10)
              kmeans.fit_predict(feature1000)
              labels = kmeans.fit_predict(feature1000)
              s[k] = metrics.silhouette_score(feature1000,labels,metric='euclidean')

          plt.plot(range(2,len(s)),s[2:])
          plt.xlabel('Number of clusters')
          plt.ylabel('Kmeans Silhouette Score')
          plt.show()
```



```
In [216]: # select k=5 to cluster
          k_K1000 = 5

          kmeans = KMeans(init='k-means++', n_clusters=k_K1000, n_init=10)
          K_labels1000 = kmeans.fit_predict(feature1000)
          K_centroids1000 = kmeans.cluster_centers_

          K_dic1000 = {}
          for i in K_labels1000:
              if i in K_dic1000:
                  K_dic1000[i] += 1
              else:
                  K_dic1000[i] = 1

          # print(K_dic10)
          for i in range(k_K1000):
              print('There are {} restautrants in cluster {}'.format(K_dic1000[i], i))
```

```
There are 1130 restautrants in cluster 0.
There are 1195 restautrants in cluster 1.
There are 2081 restautrants in cluster 2.
There are 1056 restautrants in cluster 3.
There are 437 restautrants in cluster 4.
```

```

In [217]: K_centroids1000 = kmeans.cluster_centers_
          K_ori_centers1000 = lsa.inverse_transform(K_centroids1000[:, :100])
          K_order_centroids1000 = K_ori_centers1000.argsort()[:, :-1]

          print("Top words in each cluster by using K-means method:")
          print('')

          # print out the top 50 words with largest weight in each cluster
          for i in range(k_K1000):
              print("Top 50 words in Cluster {}: ".format(i))
              for ind in K_order_centroids1000[i, :50]:
                  # print('Top category is {}'.format(terms10[ind]))
                  print(' %s' % terms1000[ind], end='')
              print("")
              print("")

```

Top words in each cluster by using K-means method:

Top 50 words in Cluster 0:

pizzas subway truck pepperoni fremont mcdonald dogs el driver pancakes hash burritos coupons

Top 50 words in Cluster 1:

filet brunch calamari patio sliders pancakes lamb asparagus steaks bartenders hash parmesan

Top 50 words in Cluster 2:

court subway pizzas tofu pepperoni curry panda korean chow teriyaki dogs mein mcdonald subs

Top 50 words in Cluster 3:

pizzas pepperoni subway teriyaki nachos bartenders phillygroupon chipotle subs games hawaiian

Top 50 words in Cluster 4:

asada carne burritos tortillas pastor guacamole carnitas enchiladas salsas nachos el quesadilla

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : pizzas, pancakes, burritos

Cluster 1 : brunch, steaks, meatballs

Cluster 2 : subways, Chinese, Korean

Cluster 3 : nachos, Mexican, Japanese

Cluster 4 : salsas, Spanish, burger

```
In [218]: K_cluster1000_top = [['pizzas, pancakes, burritos'], ['brunch, steaks, meatballs'], ['salsas, Spanish, burger'], ['subways, Chinese, Korean'], ['nachos, Mexican, Japanese'], ['brunch, steaks, meatballs']]


c = ['r', 'g', 'blue', 'yellow', 'purple', 'darkseagreen', 'dimgray', 'pink', 'brown']

plt.figure(figsize=(16,10))

for i in range(k_K1000):
    plt.scatter(latitude[K_labels1000==i], longitude[K_labels1000==i], s=25, color=c[i])

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('Kmeans Clusters of top1000 reviews in each restaurants.')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```

Out[218]: Text(0,0.5,'Longitude')



12 Hierarchical clusters

```
In [219]: H_Z1000 = hierarchy.linkage(feature1000, method='ward', metric='euclidean')

s = np.zeros(max_clusters+1)
for k in range(2,max_clusters+1):
    clusters = hierarchy.fcluster(H_Z1000, k, criterion='maxclust')
    s[k] = metrics.silhouette_score(feature1000,clusters,metric='euclidean')
plt.plot(range(2,len(s)),s[2:])
plt.xlabel('Number of clusters')
plt.ylabel('Hierarchy Silhouette Score')
```

```
Out [219]: Text(0,0.5,'Hierarchy Silhouette Score')
```



```
In [220]: k_H1000 = 5
clusters1000 = hierarchy.fcluster(H_Z1000, k_H1000, criterion='maxclust')

H_dic1000 = {}
for i in clusters1000:
    if i in H_dic1000:
        H_dic1000[i] += 1
    else:
        H_dic1000[i] = 1

for i in range(1,k_H1000+1):
    print('There are {} restaurants in cluster {}'.format(H_dic1000[i], i))
```

```
There are 1046 restaurants in cluster 1.
There are 420 restaurants in cluster 2.
There are 2555 restaurants in cluster 3.
There are 538 restaurants in cluster 4.
There are 1340 restaurants in cluster 5.
```

```
In [221]: # Calculate the centroids in each cluster
cluster_H1000 = {}
sum_H1000 = []
```



```

H_centroids1000 = []
sum_HH1000 = np.zeros(102)
for i in range(k_H1000):
    cluster_H1000[i] = np.array([feature1000[j] for j in range(len(location)) if clus
    for k in range(len(cluster_H1000[i])):
        sum_HH1000 = list(map(lambda x: x[0]+x[1], zip(sum_HH1000, cluster_H1000[i]
    sum_H1000.append(sum_HH1000)
    center = [(sum_HH1000[j]/len(cluster_H1000[i])) for j in range(len(sum_HH1000))]
    H_centroids1000.append(center)

H_centroids1000 = np.array(H_centroids1000)
H_ori_centers1000 = lsa.inverse_transform(H_centroids1000[:, :100])
H_order_centroids1000 = H_ori_centers1000.argsort()[:, :-1]

print("Top words in each cluster by using Hierarchical method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_H1000):
    print("Top 50 words in Cluster {}: ".format(i+1))
    for ind in H_order_centroids1000[i, :50]:
        print(' %s' % terms1000[ind], end='')
    print("")
    print("")

```

Top words in each cluster by using Hierarchical method:

Top 50 words in Cluster 1:

tofu curry mein chow panda teriyaki japanese tempura pad korean china wonton pho mongolian mi

Top 50 words in Cluster 2:

pizzas pepperoni tofu curry mein chow panda teriyaki japanese driver tempura pad korean hut cl

Top 50 words in Cluster 3:

pizzas pepperoni curry tofu court bartenders philly lamb chow calamari panda teriyaki eggplant

Top 50 words in Cluster 4:

pizzas asada nachos pepperoni carne court curry tofu burritos bartenders guacamole chow lamb p

Top 50 words in Cluster 5:

pizzas subway nachos pepperoni asada carne court burritos curry bartenders tofu teriyaki guac

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 1 : curry, Chinese, Japanese

Cluster 2 : pizzas, tofu, Korean

Cluster 3 : pancakes, meatballs, brunch

Cluster 4 : nachos, burritos, Mexican

Cluster 5 : pizzas, subway, burgers

```
In [222]: H_cluster1000_top = [['curry, Chinese, Japanese'], ['pizzas, tofu, Korean'], ['pancakes, meatballs, brunch'], ['nachos, burritos, Mexican'], ['pizzas, subway, burgers']]
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_H100):
    plt.scatter(latitude[clusters1000==i+1], longitude[clusters1000==i+1], s=25, color=c[i])

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('Hierachy Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```

Out [222]: Text(0,0.5,'Longitude')



13 GMM clusters


```
In [179]: s = np.zeros(max_clusters+1)
for k in range(2,max_clusters+1):
    gmm = mixture.GaussianMixture(n_components=k, covariance_type='full')
    gmm.fit(feature1000)
```

```

y_pred = gmm.predict(feature1000)
s[k] = metrics.silhouette_score(feature1000,y_pred,metric='euclidean')
plt.plot(range(2,len(s)),s[2:])
plt.xlabel('Number of clusters')
plt.ylabel('GMM Silhouette Score')

```

Out[179]: Text(0,0.5,'GMM Silhouette Score')



```

In [223]: k_G1000 = 5
gmm = mixture.GaussianMixture(n_components=k_G1000, covariance_type='tied')
gmm.fit(feature1000)
y_pred1000 = gmm.predict(feature1000)

G_dic1000 = {}
for i in y_pred1000:
    if i in G_dic1000:
        G_dic1000[i] += 1
    else:
        G_dic1000[i] = 1

for i in range(k_G1000):
    print('There are {} restaurants in cluster {}'.format(G_dic1000[i], i))

```

There are 531 restaurants in cluster 0.
There are 668 restaurants in cluster 1.

There are 471 restaurants in cluster 2.
 There are 3663 restaurants in cluster 3.
 There are 566 restaurants in cluster 4.

```
In [224]: # Calculate the centroids in each cluster
cluster_G1000 = {}
sum_G1000 = []
G_centroids1000 = []
sum_GG1000 = np.zeros(102)
for i in range(k_G1000):
    cluster_G1000[i] = np.array([feature1000[j] for j in range(len(location)) if y_p
    for k in range(len(cluster_G1000[i])):
        sum_GG1000 = list(map(lambda x: x[0]+x[1], zip(sum_GG1000, cluster_G1000[i]
    sum_G1000.append(sum_GG1000)
    center = [(sum_GG1000[j]/len(cluster_G1000[i])) for j in range(len(sum_GG1000))]
    G_centroids1000.append(center)

G_centroids1000 = np.array(G_centroids1000)
G_ori_centers1000 = lsa.inverse_transform(G_centroids1000[:, :100])
G_order_centroids1000 = G_ori_centers1000.argsort()[:, :-1]

print("Top words in each cluster by using GMM method:")
print('')

# print out the top 50 words with largest weight in each cluster
for i in range(k_G1000):
    print("Top 50 words in Cluster {}: ".format(i))
    for ind in G_order_centroids1000[i, :50]:
        print(' %s' % terms1000[ind], end='')
    print("")
    print("")
```

Top words in each cluster by using GMM method:

Top 50 words in Cluster 0:

filet steakhouse steaks mignon calamari ravioli veal asparagus ribeye lamb olive scallops spa

Top 50 words in Cluster 1:

bartenders nachos machines burritos poker gaming filet games chipotle steaks pt calamari vide

Top 50 words in Cluster 2:

pizzas pepperoni bartenders hut meatballs mozzarella parmesan driver marinara knots machines

Top 50 words in Cluster 3:

pizzas subway pepperoni court curry bartenders tofu nachos teriyaki mcdonaldgroupon pancakes

Top 50 words in Cluster 4:

pizzas subway nachos pepperoni asada carne court burritos curry bartenders tofu teriyaki guac

From top 50 words in each cluster, we can select several key words as their labels.

Cluster 0 : steaks, meatballs, wines

Cluster 1 : nachos, Mexican, Hawaiian

Cluster 2 : pizzas, burritos, spaghetti

Cluster 3 : subway, curry, burgers

Cluster 4 : carne, Korean, Japanese

```
In [225]: G_cluster1000_top = [['steaks, meatballs, wines'], ['nachos, Mexican, Hawaiian'], ['pizzas, burritos, spaghetti'], ['subway, curry, burgers'], ['carne, Korean, Japanese']]
c = ['r', 'g', 'blue', 'yellow', 'burlywood', 'darkseagreen', 'beige', 'pink', 'orange']

plt.figure(figsize=(16,10))
for i in range(k_G1000):
    plt.scatter(latitude[y_pred1000==i], longitude[y_pred1000==i], s=25, color=c[i], label=G_cluster1000_top[i])

plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), prop={'size': 15})
plt.title('GMM Clusters')
plt.xlabel('Latitude')
plt.ylabel('Longitude')
```

Out[225]: Text(0,0.5,'Longitude')



What observations can you make regarding the different k values and clusterings? We are expecting comments on the labels of each cluster for the different k , on the districts created, on the results of each clustering e.t.c. In general, feel free to report any interesting findings you made. (2 pts)

13.0.1 Findings

1. It took me nearly 5 hours to run stemming of 100 or 1000 top reviews. So I gave up, only use the words which were not stemmed... I am so sorry about it.
 2. When $k=10$, by using the stemmed information, we can get more precise high frequency key words in the all reviews because we eliminate the noises from the words endings (That's why "Mexican" frequency is much less in $k=100$ and $k=1000$ than $k=10$)
 3. When k becomes larger, the distribution of clusters weight becomes more even.
 4. From the scatter figures, the GMM method has most overlapping clusters.
 5. Key words about Fast food tend to be high frequency key words in each clusters.
-